

ATHARVA JOSHI

Salt Lake City, UT | (+1) 763-202-0270 | [Email](#) | [Linkedin](#) | [Github](#) | [Medium](#)

SUMMARY

GenAI engineer fluent in development of RAG systems, AI agents, MCPs, support copilots, and NLP/forecasting engines. Expert in designing and implementing AI/ML solutions with a strong developer mindset for sellers, analysts, and strategists in retail, pricing analytics, and martech spaces. 4+ years building systems with proven scalability and enterprise adoption.

WORK EXPERIENCE

ALTIMETRIK CORP.

Gen AI Engineer

San Francisco, CA

Nov 2024 - Present

Lululemon - AI Powered Workflow Automation Solution for Feature Change Requests (LangGraph, LangSmith, OpenAI LLMs, AWS Bedrock, ChromaDB, GitHub MCP)

- Spearheaded architecting of an org-wide first-of-its-kind GenAI-powered change request automation framework to **reduce manual effort across analysts, engineers, and QE** while establishing a repeatable blueprint for **data workflow governance** using LangGraph to orchestrate LLM agents across parsing, impact analysis, SQL generation, and validation.
- Implemented **dual retrieval strategy** with ChromaDB (RAG over PDF knowledge base) and Neo4j lineage graph, enabling both semantic grounding from functional design docs and structural impact analysis across tables/views and **eliminating error-prone dependencies and reduced regression risk** in schema modifications.
- Integrated OpenAI and AWS Bedrock foundational models with LangSmith for tracing and observability, ensuring outputs were **auditable, testable, and production-grade**. Delivered validated SQL bundles and rollback plans that **reduced request-to-implementation cycle time by 60-70%**.
- Led cross-functional adoption by exposing the pipeline via GitHub MCP and TypeScript UI, empowering business analysts to self-serve schema changes while maintaining engineering oversight — **cutting turnaround from weeks to days** and freeing senior engineers for higher-value initiatives.

Customer Meeting Prep Agent for an Enterprise Monitoring Platform (OpenAI Agents SDK, Pinecone, AWS, GitHub Actions)

- Architected and deployed a modular ingestion pipeline integrating 5+ data sources to vectorize 100K+ user story fragments, **powering a context-aware GenAI assistant** for generating pre-meeting insights.
- Enabled sales teams to **reduce prep time** by 45–50 minutes per meeting, enhancing productivity and increasing deal win rates.
- Engineered agentic workflow using OpenAI's Agents SDK, dynamically fetching and synthesizing insights from Pinecone Vector DB tailored to meeting context—boosting seller confidence and **call personalization across the organization**.
- Documented and standardized AWS-based deployment process using GitHub Actions, building a reusable enterprise template that **improved implementation velocity** by 60% and enabled **parallel adoption** by adjacent teams.
- Mentored 2 junior engineers in productionizing the pipeline, scaling knowledge across teams and establishing operational resilience.

Customer Support Assistance for a Prominent Pricing Analytics Company (AWS, Pinecone, OpenAI Custom GPTs)

- Leveraged OpenAI embedding models and Pinecone to **create a vectorized knowledge base** by structuring and ingesting approx. 100GB of historical support data spanning multiple formats.
- **Reduced manual support search time** by 75% by delivering a RAG-based CustomGPT assistant, enabling **faster case resolution**.

Product Performance Analysis for a Leading American Women's Apparel Brand (Python, SQL, Snowflake Cortex)

- Unlocked approx. \$135K in potential revenue by **identifying hidden customer segments** using LLaMA from Snowflake Cortex to analyze multi-channel product reviews.

CARLSON ANALYTICS LAB

Analytics Lead

Minneapolis, MN

Jul 2023 - May 2024

Impact Analysis for Insiders Loyalty Program for a major Minneapolis mall. (Python, SQL, PowerBI, EDA, JIRA)

- Led cross-functional analytics initiative aligning marketing, customer engagement, and strategy teams to enhance loyalty program ROI through data-driven decision-making.
- Spearheaded **customer segmentation** of 13K+ insiders using K-Means clustering, increasing targeting precision and campaign efficiency by 30%.
- Unlocked 23% surge in engagement by designing **personalized promotional strategies** using association rule mining.

ALTIMETRIK INDIA

Associate Data Scientist

Bangalore, India

Nov 2020 - May 2023

Clients: Leading Japanese sportswear & apparel company and a leading Canadian athleisure designer & manufacturer. (PyCaret, RNN, ML explainability, NLP)

- Delivered \$150K in cost savings by optimizing **SKU planning** using ML pipelines. Boosted quarterly profits by 18% via **seasonal product forecasting** with LSTM models.
- Reduced **inventory re-stock delays** by 70% (from 7 to 2 days) using Python-based time series models. Saved \$70.5K through BERT-based **employee sentiment classification**.

Client: CleverTap – US customer lifecycle management & marketing services company (BigQuery, CI/CD, Spark)

- Engineered **churn prediction** models using decision trees that retained \$150K in annual revenue. Built automated **customer journey pipelines** in BigQuery, realizing \$400K in savings annually through faster analysis of 14M+ monthly records.

TECHNICAL SKILLS PORTFOLIO

Strategic Technology Leadership

- GenAI Systems Architecture | Enterprise AI Strategy & Deployment
- ML Lifecycle Ownership | Cloud-based MLOps (AWS, GitHub Actions)
- Cross-functional Enablement | Knowledge Scaling | Technical Mentorship

Advanced AI, ML & Analytics

- Retrieval-Augmented Generation (RAG), Prompt Engineering, Agentic AI Design
- LLM Orchestration: OpenAI Agents SDK, LLaMA, BERT, Custom GPTs
- Predictive Modeling, Demand Forecasting, NLP, Churn Prediction
- Machine Learning: LSTM, RNN, Decision Trees, PyCaret, Model Explainability

Enterprise Infrastructure & Platforms

- Vector DBs: Pinecone — Data Lakes & Pipelines: S3, Snowflake, BigQuery
- Cloud: AWS (Lambda, EC2), GCP, Snowflake Cortex — CI/CD: GitHub Actions, Airflow
- Databases and Warehouse: Snowflake, MySQL, PostgreSQL
- IDEs and Editors: Cursor AI, Trae, VS Code, PyCharm

Domain & Business Expertise

- Industry Focus: SaaS, Retail & Apparel, Martech, Loyalty Programs
- Business Applications: CX Optimization, Sales Enablement, Operational Forecasting
- Strategic Insight Generation — Enterprise Data Activation

EDUCATION

CARLSON SCHOOL OF MANAGEMENT, UNIVERSITY OF MINNESOTA

May 2024

Master of Science in Business Analytics

University of Pune, India

May 2020

B.Eng, Computer Engineering